Research brief | AITR 7

Chinmay Thakur | MDE2026011 | IIIT Allahabad | 18 October 2026

Review and sources	T23: citation integrity, with presentation, identity, access and claim support separated. Verified original sources: Abbonato (2026), Badalova & Mayr (2026), Shi et al. (2026), and Gao et al. (2023).
Research gap	The released checker decisions permit testing whether simple agreement reduces false alarms. This is a tractable inferred question, not a claim of a new voting algorithm or a systematic novelty result.
Primary question	Does a fixed three-of-five majority reduce RefChecker's false-positive rate by at least 0.10 absolute while retaining recall of at least 0.80 on the same corpus?
Supporting questions	How do thresholds, panel removals and source-held-out selection change the trade-off? What review workload and residual errors follow from fixed disagreement deferral?
Pathway / build	Controlled comparative reanalysis. Implement a strict data loader, pure voting/triage functions, grouped threshold selection, metrics, tests and reproducible evidence exports. Original services are not rerun.
Data and versions	104 references from three source documents, five checker flags each. Zenodo v1.0.0: 33 problematic / 71 verified. Author-corrected v1.1.0: 35 / 69. Both CC BY 4.0; raw bytes, metadata and hashes retained.
Comparators / controls	All five released checkers; k=1..5 thresholds; leave-one-document-out selection; five strict 3-of-4 removals. Same rows and labels within a release. Labels never enter fixed predictors; held-out labels never choose their own threshold.
Measures / depth	Primary recall, false-positive rate and balanced accuracy, with confusion counts. Threshold and panel sensitivity, per-document analysis, triage coverage, error patterns, and conditional paired bootstrap (5,000; seed 20261018).
Feasibility / risks	CPU-only tabular processing, no models or paid APIs. Small purposive corpus, inherited labels, heterogeneous sources, conservative duplicate screen, no human-trust or semantic-support measurements.
Verified checkpoint	Historical RefChecker matrix reproduced: TP32 / FP36 / TN35 / FN1. Sixteen tests pass; 24 numerical output files match a clean-environment rerun. Corrected majority has 36 false alarms and 4 misses: the primary target fails.
Provenance	Protocol frozen before reference-level outcomes; published aggregates already known. Corrected release discovered after the first run and analyzed under amendment A2. Student review and live demonstration remain personal responsibilities.

Official data: doi.org/10.5281/zenodo.21457492 and doi.org/10.5281/zenodo.23236264. Full source links, methods, evidence and amendments are included in the repository.